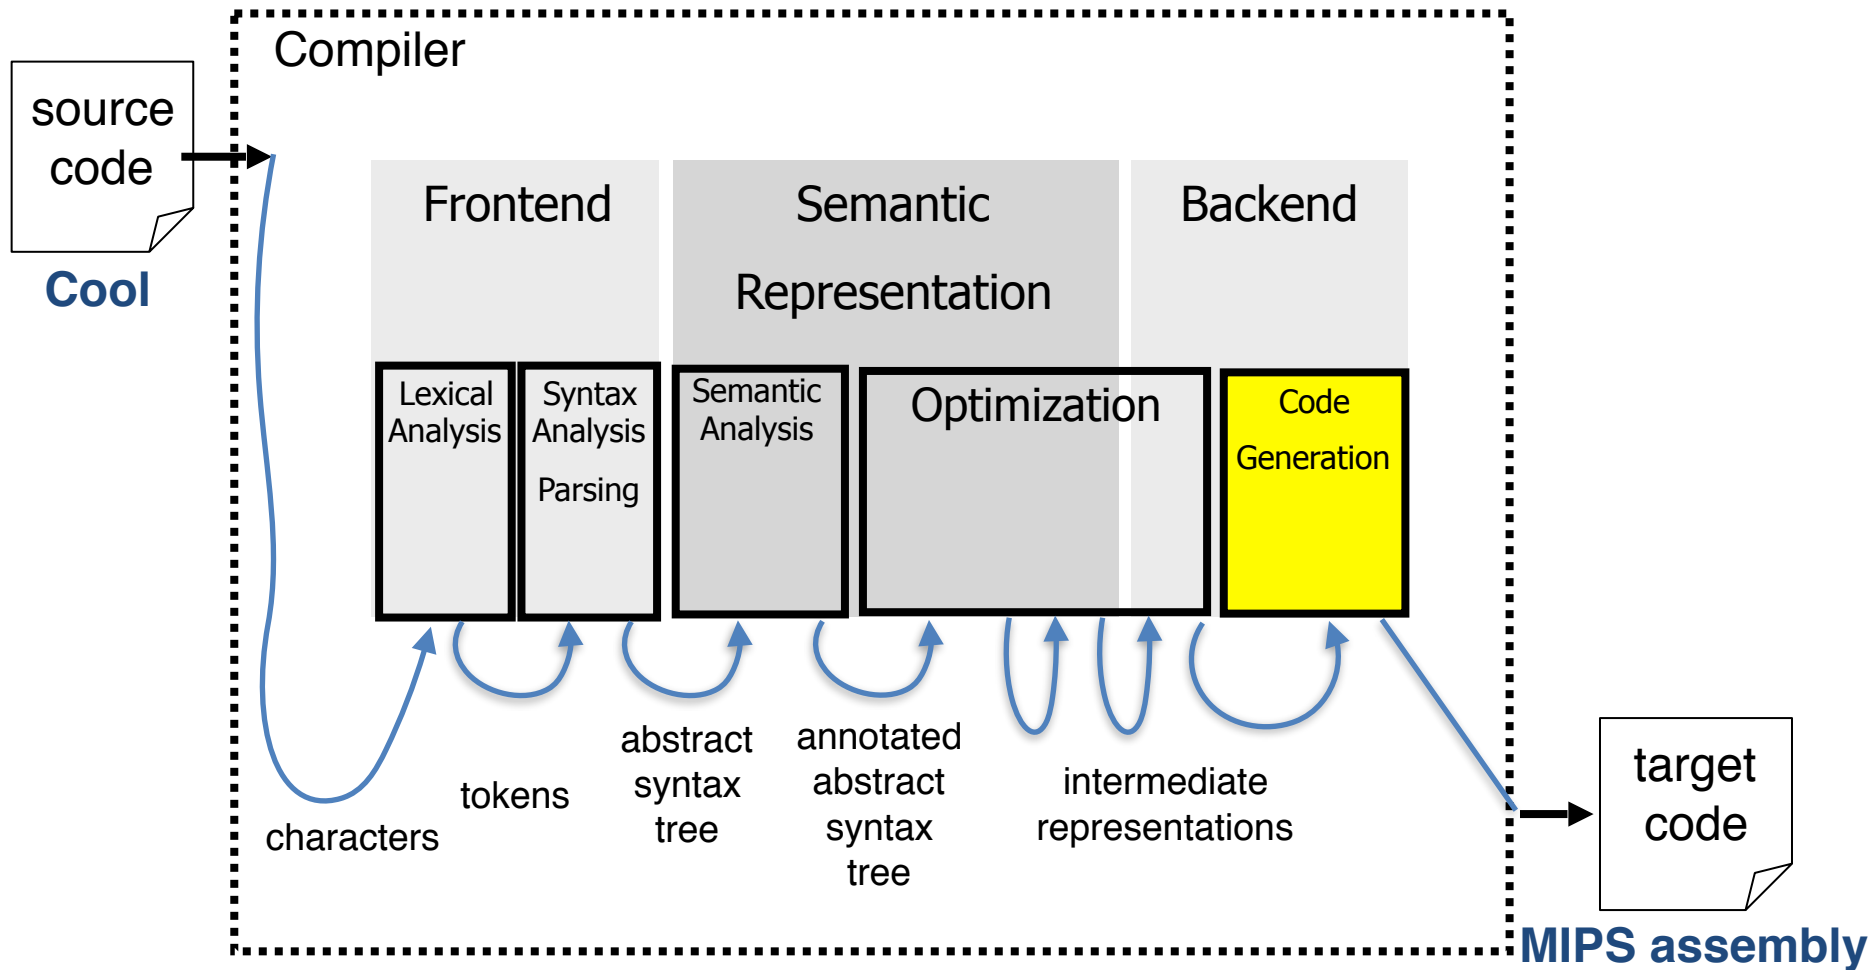


# Runtime organization


STACK FRAMES  
(ACTIVATION RECORDS)



# Supporting function calls

- How is that done?
- What do we need from the compiler?

`n = f(a[i])`



```
t1 = i * 4  
t2 = a + t1  
param t2  
t3 = call f, 1  
n = t3
```

# Supporting function calls

## ✓ Type checking

- function type: return type, type of formal parameters
- within an expression function treated like any other operator

## ✓ Symbol table

- parameter names
- New computing environment (scope)
  - at least temporary memory for local variables
- Pass information into the new environment
  - parameters
- Transfer of control to/from method and handle return values

# Design decisions

- Scoping rules
  - static scoping vs. dynamic scoping
- Caller/callee conventions
  - parameters
  - who saves register values?
- Allocating space for local variables

# Static scoping

```
main ( )
{
    int a = 0 ;
    int b = 0 ;
    {
        int b = 1 ;
        {
            int a = 2 ;
            printf ("%d %d\n", a, b);
        }
        {
            int b = 3 ;
            printf ("%d %d\n", a, b) ;
        }
        printf ("%d %d\n", a, b) ;
    }
    printf ("%d %d\n", a, b) ;
}
```

a name refers to its (closest)  
enclosing scope

known at **compile** time

Output:

2	1
0	3
0	1
0	0

# Static scoping

```
main ( )  
{  
    int a = 0 ;  
    int b = 0 ;  
    {  
        int b = 1 ;  
        {  
            B2    int a = 2 ;  
                  printf ("%d %d\n", a, b);  
        }  
        {  
            B3    int b = 3 ;  
                  printf ("%d %d\n", a, b) ;  
        }  
        printf ("%d %d\n", a, b) ;  
    }  
    printf ("%d %d\n", a, b) ;  
}
```

Diagram illustrating static scoping with nested scopes labeled B<sub>0</sub>, B<sub>1</sub>, B<sub>2</sub>, and B<sub>3</sub>. B<sub>0</sub> is the outermost scope (main), followed by B<sub>1</sub>, B<sub>2</sub>, and B<sub>3</sub> as inner scopes.

a name refers to its (closest)  
enclosing scope

known at **compile** time

Declaration	Scopes
a=0	B0,B1,B3
b=0	B0
b=1	B1,B2
a=2	B2
b=3	B3

# Dynamic scoping

- Scope determined at runtime
- Each identifier is associated with a global stack of bindings
- When entering scope where identifier is declared
  - push declaration on identifier stack
- When exiting scope where identifier is declared
  - pop identifier stack
- **Evaluating the identifier in any context binds to the current top of stack**



# Example

```
int x = 42;  
  
int f() { return x; }  
int g() { int x = 1; return f(); }  
int main() { return g(); }
```

- What value is returned from main?
  - static scoping
  - dynamic scoping

# Example: static vs dynamic scope

```
int x = 37;  
int y = 42;  
void Function1() {  
    Print(x + y);  
}  
void Function2() {  
    int x = 0;  
    Function1();  
}  
void Function3() {  
    int y = 0;  
    Function2();  
}  
Function1();  
Function2();  
Function1();  
Function3();
```

static scoping

output:

79

79

**79**

79

dynamic scoping

output:

79

42

**79**

0

# Why do we care?

- We need to generate code to access variables
- Static scoping
  - identifier binding is known at compile time
  - address of the variable is known at compile time
  - assigning addresses to variables is part of code generation
  - no runtime errors of “access to undefined variable”
  - can check types of variables

# Variable addresses for static scoping first attempt

```
int x = 42;  
  
int f() { return x; }  
int g() { int x = 1; return f(); }  
int main() { return g(); }
```

identifier	address
x (global)	0x24
x (inside g)	0x72

# Variable addresses for static scoping first attempt

```
int a [11] ;

void quicksort(int m, int n) {
    int i;
    if (n > m) {
        i = partition(m, n);
        quicksort (m, i-1) ;
        quicksort (i+1, n) ;
    }

    main() {
        ...
        quicksort (1, 9) ;
    }
}
```

What is the address of variable “i” in quicksort?

- How do we handle recursion?

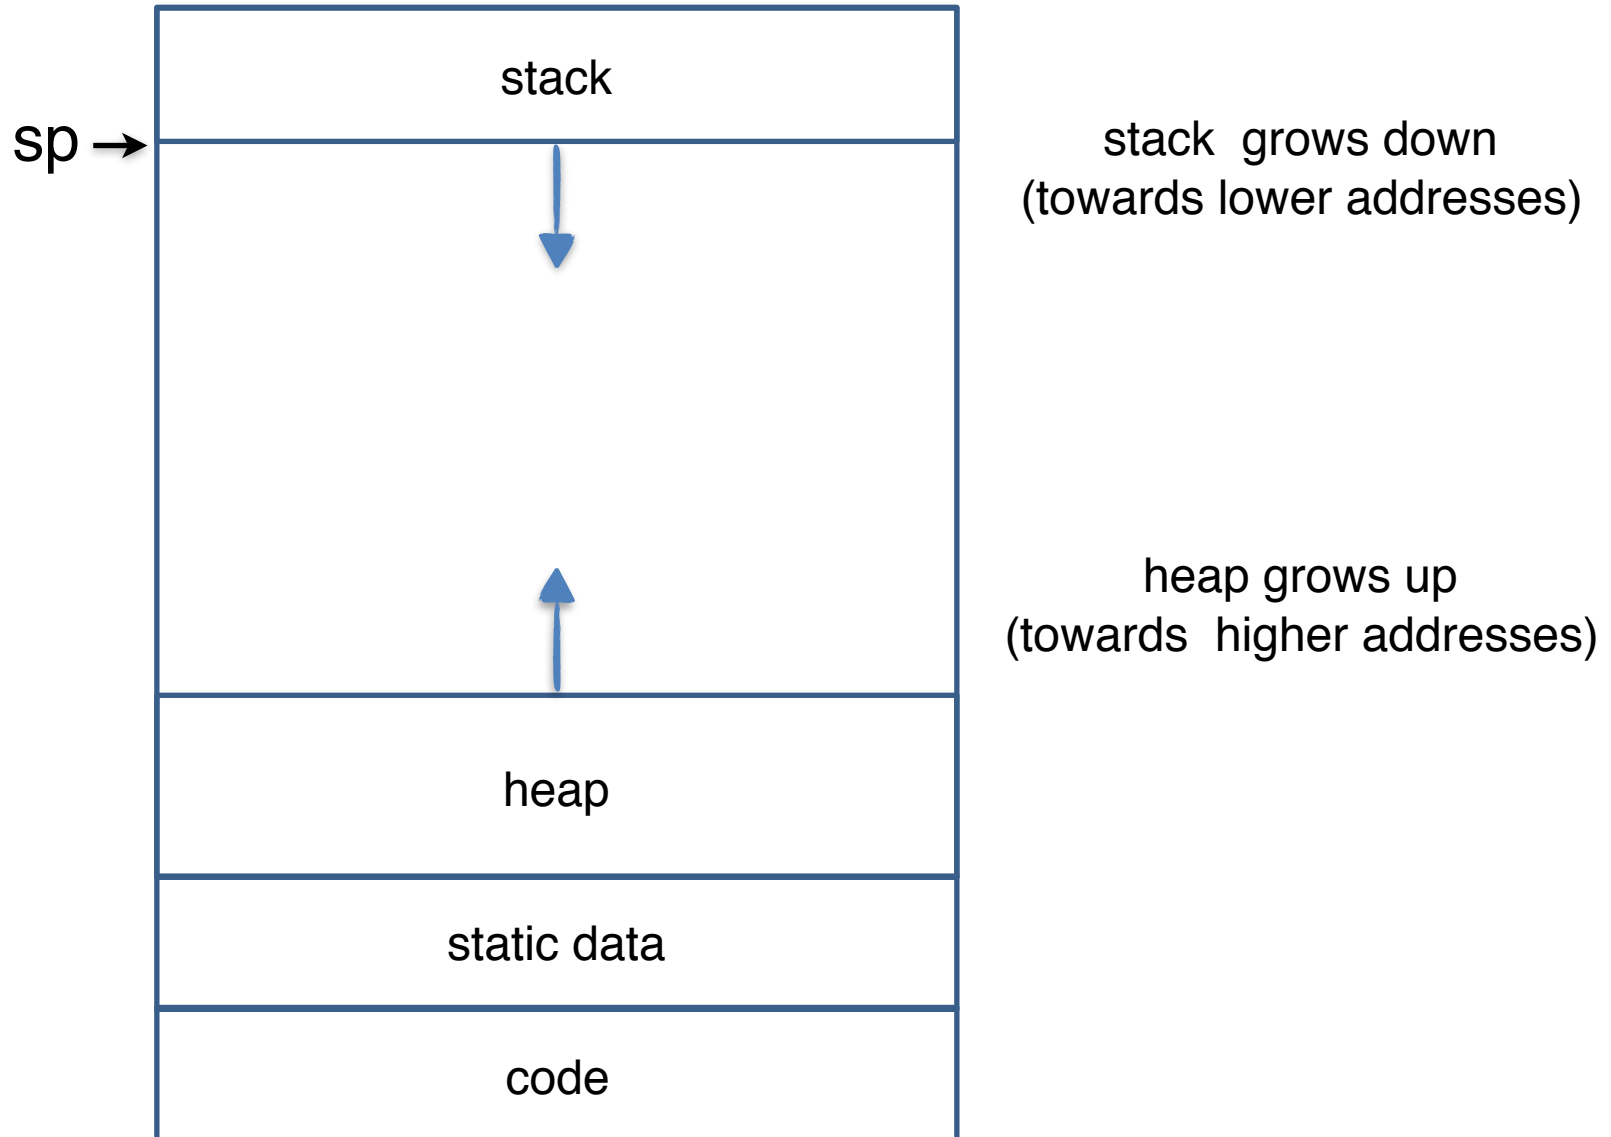
# Runtime stack

- Stack of frames
- Call: push new frame
- Return: pop frame
- Only one “active” frame: top of stack

# Stack frame

- Separate space for **each function invocation**
- Variable size
  - different functions may require different memory sizes
  - size may be input dependent
- Managed **at runtime**
  - allocated upon function call
  - deallocated upon function return
  - efficient! function are called frequently
- Code for managing it generated by the compiler


# Memory layout



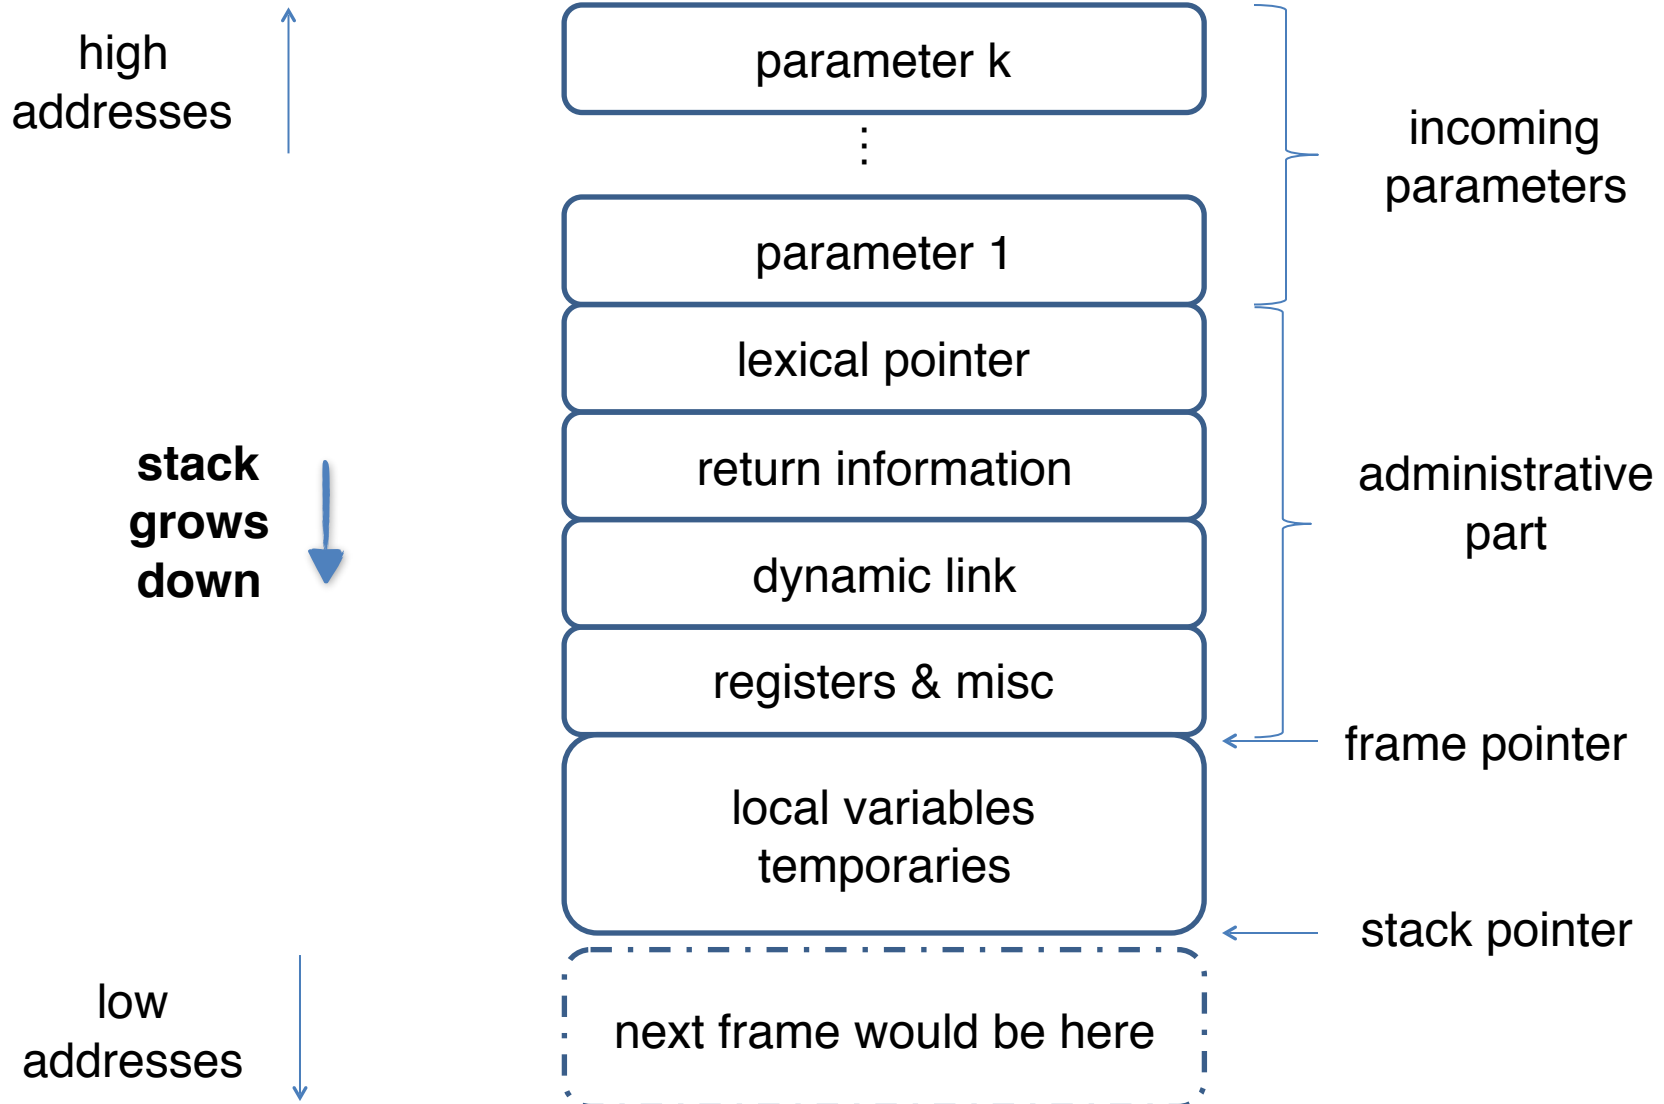


# Runtime stack

- SP is **stack pointer**
  - top of current frame
- FP is **frame pointer**
  - base of current frame
  - sometimes called BP (base pointer)



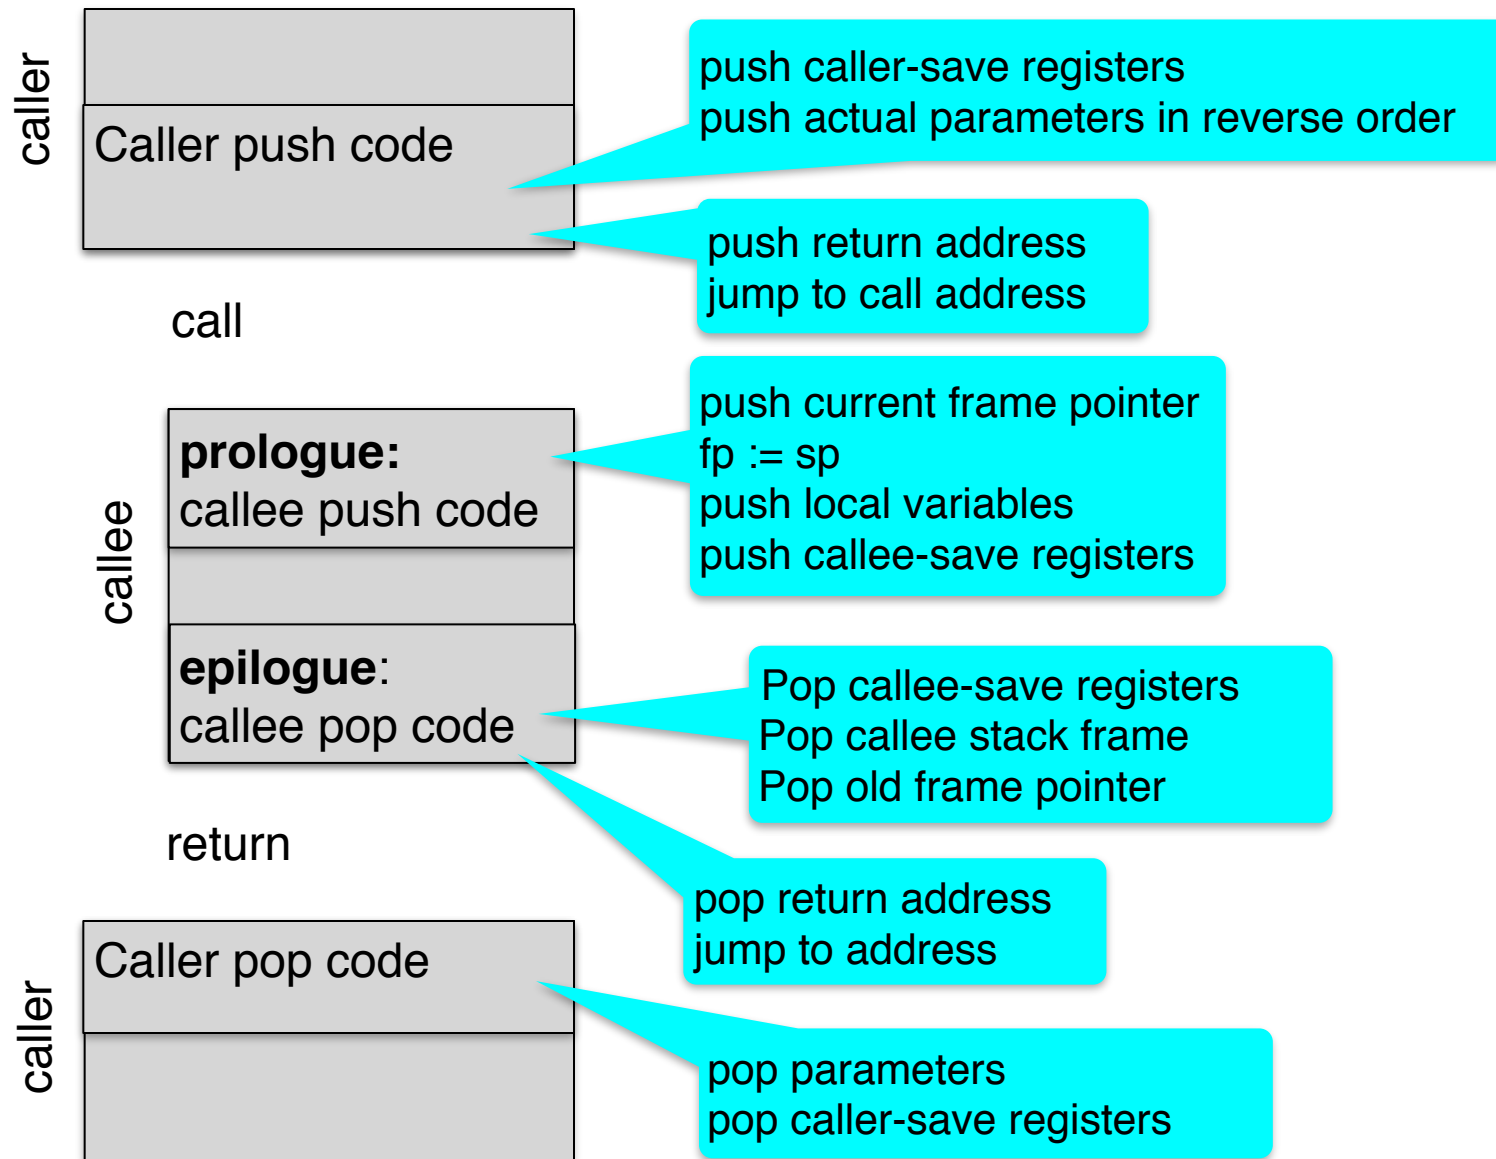
# Stack frame



# Call sequence

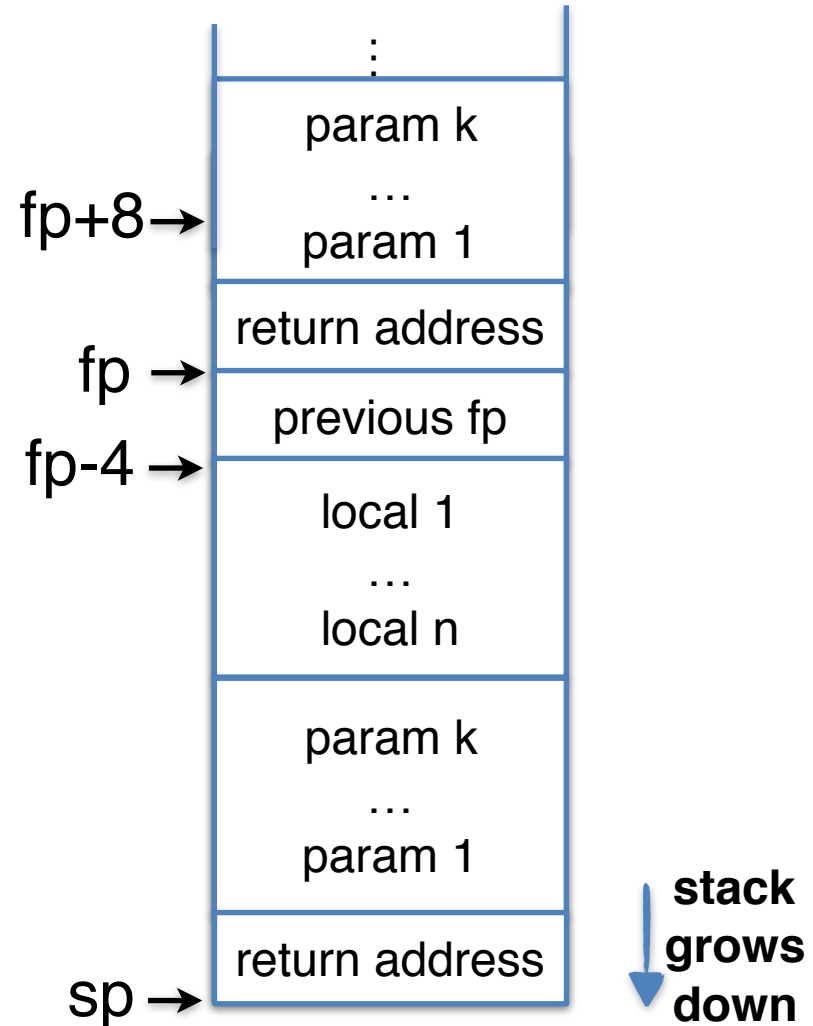
- **The processor does not save the content of registers on function calls**
- So who will?
  - caller saves and restores registers
  - callee saves and restores registers
  - but can also have each save/restore some registers

# Call sequence



# Accessing stack variables


- Use offset from fp
- Remember:  
stack grows downwards
- Above fp = parameters
- Below fp = locals
- Examples
  - $fp + 4 = \text{return address}$
  - $fp + 8 = \text{first parameter}$
  - $fp - 4 = \text{first local}$



# Buffer overflow

```
void foo (char *x) {  
    char buf[2];  
    strcpy(buf, x);  
}  
int main (int argc, char *argv[]) {  
    foo(argv[1]);  
}
```

```
gcc bf  
./a.out abracadabra  
Segmentation fault
```



# Buffer overflow attack

```
int check_authentication(char *password) {  
    int auth_flag = 0;  
    char password_buffer[16];  
  
    strcpy(password_buffer, password);  
    if(strcmp(password_buffer, "brillig") == 0)  
        auth_flag = 1;  
    if(strcmp(password_buffer, "outgrabe") == 0)  
        auth_flag = 1;  
    return auth_flag;  
}  
  
int main(int argc, char *argv[]) {    if(argc < 2) {  
    printf("Usage: %s <password>\n", argv[0]); exit(0);    }  
    if(check_authentication(argv[1])) {  
        printf("\n-----\n");  
        printf("        Access Granted.\n");  
        printf("-----\n");    }  
    else {  
        printf("\nAccess Denied.\n");  
    }  
}
```

# Stack frames

- Allocate a separate space for every function invocation
- Naturally supports recursion
- Efficient memory allocation policy
- Provides a simple way to achieve modularity



# **CALLING CONVENTIONS**

# Calling conventions: who and why?

- Microprocessor manufacturers specify “standard” schemes to be used by all compilers
  - stack layout
  - registers for parameter passing and return values
  - callee vs caller saved registers
- Functions compiled with one compiler can call functions compiled with another
- Essential for interaction with libraries and runtime system

# Parameter passing

- 1960s
  - in **memory**
  - no recursion is allowed
- 1970s
  - in **stack**
- 1980s
  - in **registers**
  - first k parameters are passed in registers (k=4 or k=6)

# Modern computer architectures

- Parameters
  - first k parameters are passed in registers
  - others on the stack
- Return address
  - automatically saved in a register on a call
  - a non-leaf function saves this value on the stack
- Function result
  - normally saved in a register on a return
- No stack support in the hardware (why?)

# Stack Operations in RISC

- PUSH
  - sub \$sp, 4
  - sw \$ra, (\$sp)
- POP
  - lw \$ra, (\$sp)
  - add \$sp, 4
- TOP
  - lw \$ra, (\$sp)

# MIPS instructions for calls

Instruction	Meaning
jal my_proc	jump and link start procedure my_proc \$ra holds address of instruction following the jal
jr \$ra	jump register return from procedure call puts \$ra value back into the PC

# MIPS Calling Convention

- Caller
  - pass arguments: first 4 are in \$a0-\$a3 the rest pushed on the stack
  - save caller-saved registers, including \$t0-\$t9 if needed
  - jal (\$ra gets address of instruction following the jal)
- Callee
  - save callee-saved registers in the frame
    - \$fp
    - \$ra (if the callee is not a leaf)
    - \$s0-\$s7 (if used by the callee)
  - push a stack frame \$fp := \$sp
  - ..... do some work
  - return value is in \$v0
  - restore callee-saved registers
  - pop the stack frame \$sp := \$fp
  - jr \$ra (puts \$ra value back into the program counter \$pc )

# Factorial Example

```
int factorial (int n){  
    if (n < 1) return 1;  
    return (n * factorial (n-1));  
}
```

```
factorial:  
    bgtz    $a0    doit  
    li      $v0    1          # base case, 0! = 1  
    jr      $ra  
  
doit:  
    addiu   $sp    $sp -8      # stack frame  
    sw      $s0    0($sp)      # will use for argument n  
    sw      $ra    4($sp)      # return address  
    move    $s0    $a0          # save argument  
    addiu   $a0    $a0 -1       # n-1  
    jal     factorial          # v0 := (n-1)!  
    mul     $v0    $s0 $v0      # n*(n-1)!  
    lw      $s0    0($sp)      # restore registers from stack  
    lw      $ra    4($sp)  
    addiu   $sp    $sp 8  
    jr      $ra
```



# To callee-save or caller-save?

- Callee-saved registers will contain the same value before and after the call
- Callee-saved registers need only be saved when callee modifies their value
- Caller-saved registers need only be saved when their value is used after call
- Callee can use caller-save registers without saving
- Caller need not save callee-saved registers before a call
- Placement of values into callee-save vs. caller-save registers determined by the **register allocator**
- Some heuristics and conventions are followed

# To callee-save or caller-save?

```
int foo(int a) {  
    int b=a+1;  
    f1();  
    g1(b);  
    return b+2;  
}
```

```
void bar (int y) {  
    int x=y+1;  
    f2(y);  
    g2(2);  
}
```

# Where is time saved?

- Most procedures are leaf procedures
- Interprocedural register allocation
- Many of the registers may be dead before another invocation
- Register windows are allocated in some architectures per call
  - Sun's Sparc

# **RUNTIME STACK: ADVANCED TOPICS**

# Nested procedures

- Pascal
  - any routine can have sub-routines
  - any sub-routine can access anything that is defined in its containing scope or inside the sub-routine itself
  - “non-local” variables
- Caml, Scala, Javascript, C#, Java8,...
- lambda expressions, closures
- Are frames allocated on stack or heap?

# Example: nested procedures

```
void p()  
{  
    int x;  
    void a()  
    {  
        int y;  
        void b () { ... c() ... };  
        void c ()  
        {  
            int z;  
            void d () { y = x + z };  
            ... b() ... d() ...  
        }  
        ... a() ... c() ...  
    }  
    a()  
}
```

$B_0$   $B_1$   $B_2$

possible call sequence:  
 $p \rightarrow a \rightarrow a \rightarrow c \rightarrow b \rightarrow c \rightarrow d$

what is the address of  
variable “y” in procedure d?

# Nested procedures

- Procedures may need to access variables of another procedure that contains it
- How do you find the right stack frame at runtime?
- When **c** uses variables from **a**, which instance of **a** to use?

# Lexical pointer

- Points to the **last frame** of the (static) nesting level above it
- Created at runtime and stored in the stack frame
- Lexical pointer of **c** points to a stack frame of **a**
- Accessing a variable requires a chain of indirect memory references through lexical pointers
- Number of links in the chain is
  - the difference in nesting depth between declaration scope and use scope of the variable
  - known at compile time




# Example: nested procedures

```
void p()  
{  
    int x;  
    void a()  
    {  
        int y;  
        void b () { ... c () ... };  
        void c ()  
        {  
            int z;  
            void d () { y = x + z };  
            ... b() ... d() ...  
        }  
        ... a() ... c() ...  
    }  
    a()  
}
```

$B_0$        $B_1$        $B_2$

possible call sequence:  
 $p \rightarrow a \rightarrow a \rightarrow c \rightarrow b \rightarrow c \rightarrow d$



# Terminology

- **old frame pointer** = dynamic link = control link  
points to the stack frame of the caller
- **lexical pointer** = static link = access link  
points to the last frame of the closest lexically  
enclosing block in program text

Is lexical pointer needed in C?

# Old frame pointer

- **Optional** if size of caller's stack frame is known at compile time
- **Convenient** in many situations:
  - fixed offsets
    - positive for locals
    - negative for parameters
  - quickly restore stack pointer during return sequence
  - dynamic allocation on stack: `alloca(...)` in C/C++
  - variable length parameter lists: `vararg` in C/C++

# Tail call elimination

- Tail call is a call performed as **the last operation** of the caller
- Optimization: set return address to that of caller
- Can we do the same with old frame pointer?
- Can we reuse the entire frame of the caller?
- Does it prevent stack overflow?
- Does it work with dynamic scope?

```
void f0 () {  
    f1(1);  
    f2(2);  
}
```

```
int f1 (int x) {  
    int y;  
    if (x > 0)  
        return f2(x-1);  
    else  
        return f2(x);  
}
```

```
int f2 (in y) {  
    return y+42;  
}
```

```
int f3 (int z) {  
    return f2(z)+42;  
}
```

# Tail recursion

- Tail recursive function is equivalent to a loop
- Example: compute the least power of 2 greater than y

```
int g(int y) { g1(1,y); }  
int g1(int x, int y) {  
    if (x>y)  
        return x;  
    else  
        return g1(2*x, y);  
}
```

```
int g(int y) {  
    int x = 1;  
    while !(x>y)  
        x = 2*x;  
    return x;  
}
```

# Address of formal parameters

- If the address of a formal parameter is taken, callee writes register to stack
- Cannot save formal parameter in the temporary area of the frame, with other callee saved registers, because all parameters must be consecutive in memory
- Space for the parameter must be allocated on the stack, even if the parameter is passed in a register
- Example: C/C++

```
void foo(int x) {  
    bar(&x);  
    return x + 1;  
}
```

```
void bar(int *a) {  
    *a = *a + 2;  
}
```

# Higher-order functions

- Functions passed as argument
- Need pointer to a frame “higher up” in stack
- Function returned as the result of function call
- Local variables of a function need to remain even after function returned
- Allocate frames on the heap, not stack