

Runtime system

Outline

- ✓ Code generation for expressions
- ✓ Code generation for methods
- ✓ Code generation for objects
- ✓ Operational semantics
- Runtime system

Runtime system

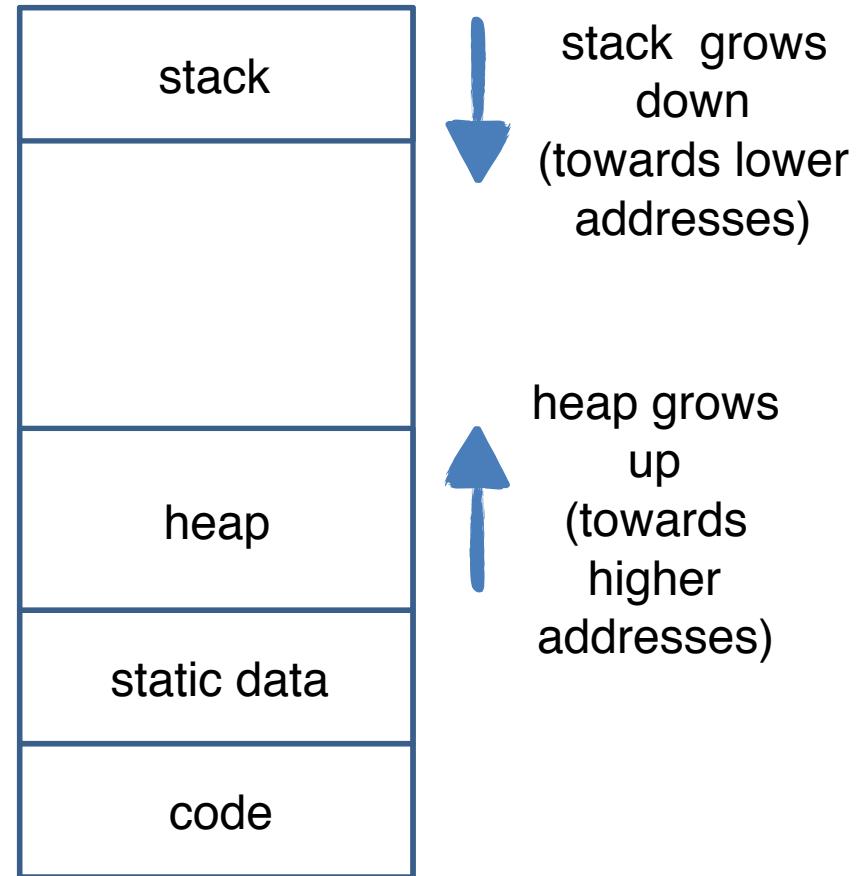
- **Mediates between OS and programming language**
- Hides details of the os and the machine from the programmer
- Ranges from simple support functions all the way to a full-fledged VM
 - Microsoft CLR for C# and other .Net languages
 - Java Virtual Machine (JVM) by Sun/Oracle or IBM for Java, Scala, ...
 - Mozilla SpiderMonkey for JavaScript
- Handles common tasks
 - memory management, including garbage collector (GC)
 - dynamic optimizations such as Just-In-Time (JIT) compiler
 - thread management
 - exception handling
 - security
 - debugging

Memory management

- Tasks
 - allocation
 - deallocation
- **Manual** memory management
 - programmer responsible for calling allocation and deallocation explicitly
- **Automatic** memory management
 - garbage collector inside runtime system automatically deallocates memory

Where do we allocate data?

- Runtime stack
 - stack frame deallocated (popped) upon method return
 - lifetime of allocated data is limited by method lifetime
- Dynamic memory allocation on the heap



Alignment

- Typically, memory access is **word-aligned**: address is a multiple of 4-bytes or 8-bytes
- Unaligned access is often slower

Alignment

- Typically, memory access is **word-aligned**: address is a multiple of 4-bytes or 8-bytes
- Unaligned access is often slower
- How do we allocate data of size 5 bytes?

Alignment

- Typically, memory access is **word-aligned**: address is a multiple of 4-bytes or 8-bytes
- Unaligned access is often slower
- How do we allocate data of size 5 bytes?
- Padding: the space until the next aligned addresses is kept empty

Manual memory management

- Examples: C, C++, Pascal, Modula, Rust

```
a = malloc(n) ;  
// do something with a  
free(a);
```

Allocating memory

```
void *malloc(size_t size)
```

- Why does malloc return void* ?

Allocating memory

```
void *malloc(size_t size)
```

- Why does malloc return void* ?
 - allocates a chunk of memory, without regard to its type

Allocating memory

```
void *malloc(size_t size)
```

- Why does malloc return void* ?
 - allocates a chunk of memory, without regard to its type
- Where is malloc implemented?

Allocating memory

```
void *malloc(size_t size)
```

- Why does malloc return void* ?
 - allocates a chunk of memory, without regard to its type
- Where is malloc implemented?
- How does it work?

Allocating memory

```
void *malloc(size_t size)
```

- Why does malloc return void* ?
 - allocates a chunk of memory, without regard to its type
- Where is malloc implemented?
- How does it work?
- How does malloc guarantee alignment?
 - after all, you don't know what type it is allocating for

Allocating memory

```
void *malloc(size_t size)
```

- Why does malloc return void* ?
 - allocates a chunk of memory, without regard to its type
- Where is malloc implemented?
- How does it work?
- How does malloc guarantee alignment?
 - after all, you don't know what type it is allocating for
 - it has to align for the largest primitive type

Allocating memory

```
void *malloc(size_t size)
```

- Why does malloc return void* ?
 - allocates a chunk of memory, without regard to its type
- Where is malloc implemented?
- How does it work?
- How does malloc guarantee alignment?
 - after all, you don't know what type it is allocating for
 - it has to align for the largest primitive type
 - in practice optimized for 8 byte alignment (glibc-2.25)

Deallocating memory

- Free too late: waste memory (memory leak)
- Free too early: dangling pointers / crashes
- Free twice: error

When can we free an object?

```
a = malloc(...) ;  
b = a;  
free (a); // ??????  
...
```

When can we free an object?

```
a = malloc(...) ;  
b = a;  
free (a); // ?????
```

When can we free an object?

```
a = malloc(...) ;  
b = a;  
free (a); // ??????  
*b = c;
```

When can we free an object?

```
a = malloc(...) ;  
b = a;  
free (a); // ??????  
*b = c;
```

```
a = malloc(...) ;  
b = a;  
free (a); // ??????  
c = malloc (...);  
if (b == c)  
    printf("unexpected equality");
```


When can we free an object?

```
a = malloc(...) ;  
b = a;  
free (a); // ??????  
*b = c;
```


```
a = malloc(...) ;  
b = a;  
free (a); // ??????  
c = malloc (...);  
if (b == c)  
    printf("unexpected equality");
```

Cannot free an object if there is a pointer to it with a future use!


When can **free x** be inserted after **p**?



When can **free x** be inserted after **p**?



When can **free x** be inserted after **p**?



on **all execution paths** after program point **p**
there are **no uses** of pointers to the object pointed to by **x**



inserting **free x** after **p** is safe

Automatic memory management

- Prevalent in object oriented languages and functional languages
- Garbage collection: automatically free memory when it is no longer needed
- Garbage collector (GC) is triggered by allocation

```
New(A) {  
  if free_list is empty then gc()  
  if free_list is empty then error("out of memory");  
  pointer := allocate(A);  
  return pointer  
}
```

Garbage collection

- Approximate reasoning about object liveness
- Use **reachability to approximate liveness**
- Assume reachable objects are live and unreachable objects are dead


Garbage collection: classical techniques

- Reference counting
- Tracing: mark and sweep
- Copying

Reference Counting GC

- Add a reference-count field to every object
 - **$o.RC$** how many references point to object o
- Newly allocated object o gets **$o.RC=1$**
- When **$o.RC=0$** the object o is unreachable
 - unreachable implies dead
 - can be collected (deallocated)

Reference Counting GC

- Add a reference-count field to every object
 - **`o.RC`** how many references point to object `o`
- Newly allocated object `o` gets **`o.RC=1`**  Why?
- When **`o.RC=0`** the object `o` is unreachable
 - unreachable implies dead
 - can be collected (deallocated)


Write-barrier for reference updates

- `collect(old)` decrement RC for all children and recursively collect objects whose RC reached 0


```
update(x,old,new) {  
    old.RC--  
    new.RC++  
    if (old.RC=0) collect(old)  
}
```

```
collect(o) {  
    free(o)  
    for c in Children(o) {  
        c.RC--;  
        if (c.RC=0) collect(c)  
    }  
}
```

Example: reference counting




Example: reference counting




`y := null`


Example: reference counting




$y := \text{null}$



Example: reference counting




$y := \text{null}$




$x := \text{null}$


Example: reference counting




$y := \text{null}$




$x := \text{null}$




Example: reference counting




$y := \text{null}$




$x := \text{null}$




Example: reference counting




$y := \text{null}$




$x := \text{null}$



Example: reference counting




$y := \text{null}$




$x := \text{null}$



Example: reference counting




$y := \text{null}$




$x := \text{null}$



Example: reference counting




$y := \text{null}$




$x := \text{null}$

y
 x

Example: reference counting




Example: reference counting




z := null

Example: reference counting



z := null



Cycles!

- Cannot identify non-reachable cycles
- Reference counts for nodes on the cycle will never decrement to 0
- Several approaches for dealing with cycles
 - ignore
 - periodically invoke a tracing algorithm to collect cycles
 - specialized algorithms for collecting cycles


Mark and Sweep GC

[McCarthy 1960]

- Marking phase
 - mark roots
 - trace all objects transitively reachable from roots
 - mark every traversed object
- Sweep phase
 - scan all objects in the heap
 - collect all unmarked objects


Mark and Sweep GC

- **Roots:** pointers in **registers** and on **stack**
- Traverse live objects and mark black
- Reclaim white objects




Mark and Sweep GC

- **Roots:** pointers in **registers** and on **stack**
- Traverse live objects and mark black
- Reclaim white objects




Mark and Sweep GC

- **Roots:** pointers in **registers** and on **stack**
- Traverse live objects and mark black
- Reclaim white objects




Mark and Sweep GC

- **Roots:** pointers in **registers** and on **stack**
- Traverse live objects and mark black
- Reclaim white objects




Mark and Sweep GC

- **Roots:** pointers in **registers** and on **stack**
- Traverse live objects and mark black
- Reclaim white objects




Mark and Sweep GC

- **Roots:** pointers in **registers** and on **stack**
- Traverse live objects and mark black
- Reclaim white objects



Mark and Sweep GC

- **Roots:** pointers in **registers** and on **stack**
- Traverse live objects and mark black
- Reclaim white objects



Mark and Sweep GC

```
mark_sweep_gc () {  
  for p in Roots  
    mark(p)  
    sweep()  
}
```

Mark and Sweep GC

```
mark (o) {  
  if (mark_bit(o) = unmarked) {  
    mark_bit(o) := marked  
    for c in Children(o)  
      mark(c)  
  }  
}
```

```
mark_sweep_gc () {  
  for p in Roots  
    mark(p)  
  sweep()  
}
```


Mark and Sweep GC

```
mark (o) {  
  if (mark_bit(o) = unmarked) {  
    mark_bit(o) := marked  
    for c in Children(o)  
      mark(c)  
  }  
}
```


```
mark_sweep_gc () {  
  for p in Roots  
    mark(p)  
  sweep()  
}
```

```
sweep() {  
  p := Heap_bottom  
  while (p < Heap_top) {  
    if (mark_bit(p) = unmarked) free(p)  
    else mark_bit(p) := unmarked;  
    p := p + size(p)  
  }  
}
```



Example: Mark and Sweep




Example: Mark and Sweep



Example: Mark and Sweep



Example: Mark and Sweep



Properties of Mark Phase

```
mark (o) {  
  if (mark_bit(o) = unmarked) {  
    mark_bit(o) :=marked  
    for c in Children(o)  
      mark(c)  
  }  
}
```

- How much memory overhead per object?

Properties of Mark Phase

```
mark (o) {  
  if (mark_bit(o) = unmarked) {  
    mark_bit(o) := marked  
    for c in Children(o)  
      mark(c)  
  }  
}
```

- How much memory overhead per object?
- Recursion depth?

Properties of Mark Phase

```
mark (o) {  
    if (mark_bit(o) = unmarked) {  
        mark_bit(o) :=marked  
        for c in Children(o)  
            mark(c)  
    }  
}
```

- How much memory overhead per object?
- Recursion depth?
- Can we traverse the heap without worst-case $O(n)$ stack?

Properties of Mark Phase

```
mark (o) {  
    if (mark_bit(o) = unmarked) {  
        mark_bit(o) := marked  
        for c in Children(o)  
            mark(c)  
    }  
}
```

- How much memory overhead per object?
- Recursion depth?
- Can we traverse the heap without worst-case $O(n)$ stack?
- Deutch-Schorr-Waite algorithm for graph marking without recursion or stack (works by reversing pointers)


Properties of Mark and Sweep GC

- Most popular method today
- Simple
- Does not move objects, so heap may **fragment**
- Complexity
 - mark phase: live objects
 - sweep phase: heap size
- Termination: each pointer traversed once
- Engineering tricks used to improve performance

Mark-Compact

- At runtime, objects are allocated and reclaimed
- Gradually, the **heap** gets **fragmented**
- When space becomes too fragmented to allocate, run **compaction algorithm**
 - **move all live objects to the beginning of the heap**
 - **update all pointers to reference the new locations**


The
Heap



Mark-Compact

- At runtime, objects are allocated and reclaimed
- Gradually, the **heap** gets **fragmented**
- When space becomes too fragmented to allocate, run **compaction algorithm**
 - **move all live objects to the beginning of the heap**
 - **update all pointers to reference the new locations**


The
Heap



Mark-Compact


- Compaction is very costly and we attempt to run it infrequently, or only partially
- Important parameters of compaction algorithm
 - keep order of objects?
 - use extra space for compactor data structures?
 - how many heap passes?
 - preserve alignment?
 - can it run in parallel on a multi-processor?

Parallel Mark and Sweep GC




Parallel GC: mutator is stopped, collector threads run in parallel


Parallel Mark and Sweep GC




Parallel GC: mutator is stopped, collector threads run in parallel



29




Concurrent Mark and Sweep GC




**Concurrent GC: mutator and collector threads run in parallel
no need to stop mutator (after roots marked)**

Concurrent Mark and Sweep GC




**Concurrent GC: mutator and collector threads run in parallel
no need to stop mutator (after roots marked)**

Concurrent Mark and Sweep GC




**Concurrent GC: mutator and collector threads run in parallel
no need to stop mutator (after roots marked)**

Concurrent Mark and Sweep GC




**Concurrent GC: mutator and collector threads run in parallel
no need to stop mutator (after roots marked)**

Concurrent Mark and Sweep GC




**Concurrent GC: mutator and collector threads run in parallel
no need to stop mutator (after roots marked)**

Concurrent Mark and Sweep GC




**Concurrent GC: mutator and collector threads run in parallel
no need to stop mutator (after roots marked)**

Concurrent Mark and Sweep GC




**Concurrent GC: mutator and collector threads run in parallel
no need to stop mutator (after roots marked)**

Concurrent Mark and Sweep GC




**Concurrent GC: mutator and collector threads run in parallel
no need to stop mutator (after roots marked)**

Concurrent Mark and Sweep GC




**Concurrent GC: mutator and collector threads run in parallel
no need to stop mutator (after roots marked)**

Concurrent Mark and Sweep GC



**Concurrent GC: mutator and collector threads run in parallel
no need to stop mutator (after roots marked)**


Concurrent Mark and Sweep GC



Concurrent GC: mutator and collector threads run in parallel no need to stop mutator (after roots marked)

Concurrent GC: interference

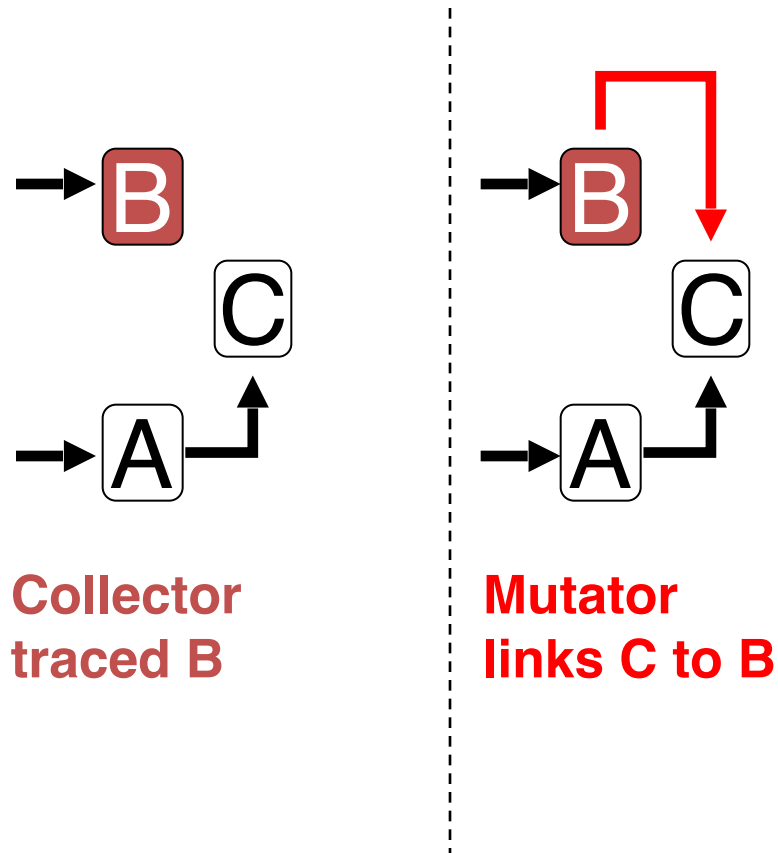
SYSTEM = MUTATOR II COLLECTOR



**Collector
traced B**

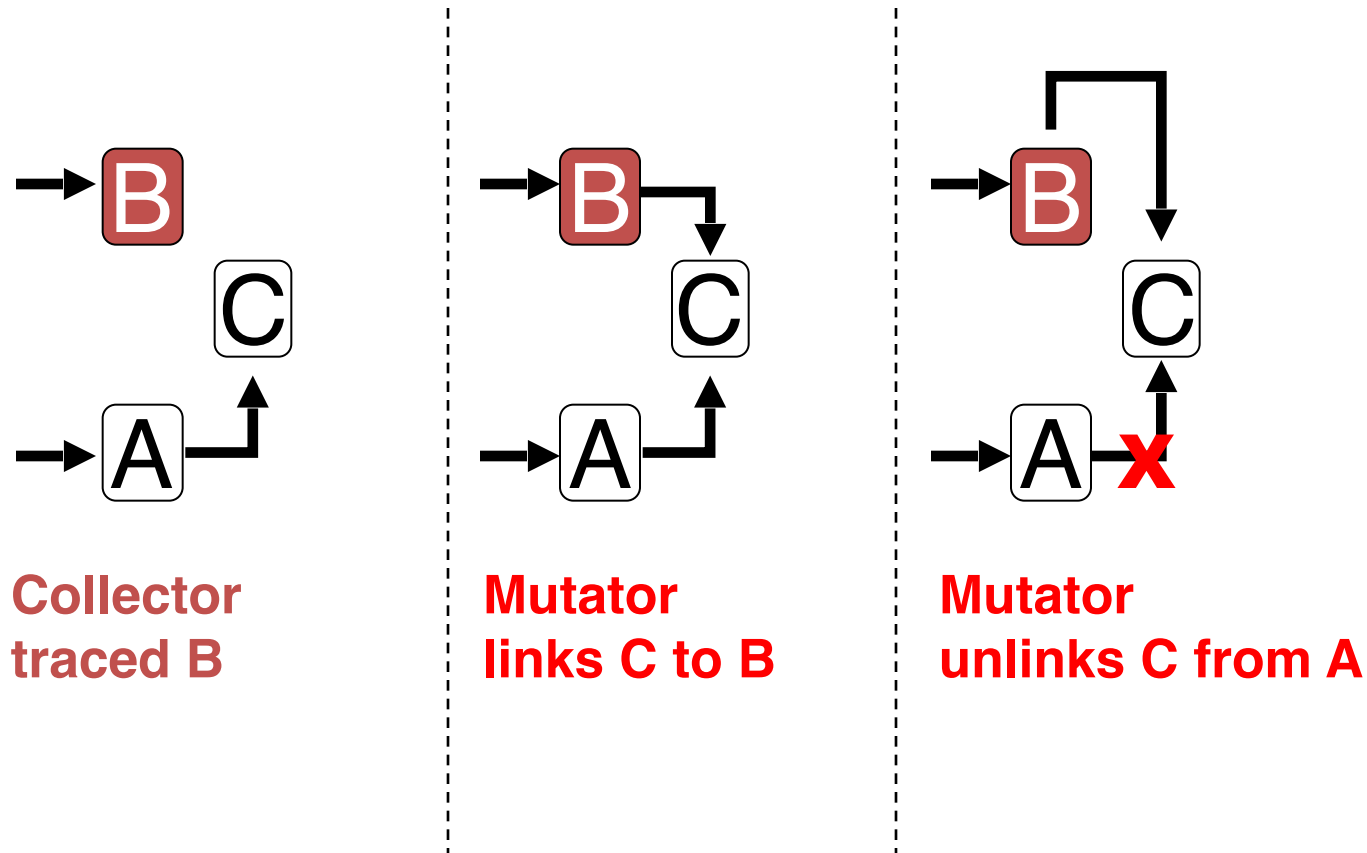
Concurrent GC: interference

SYSTEM = MUTATOR II COLLECTOR



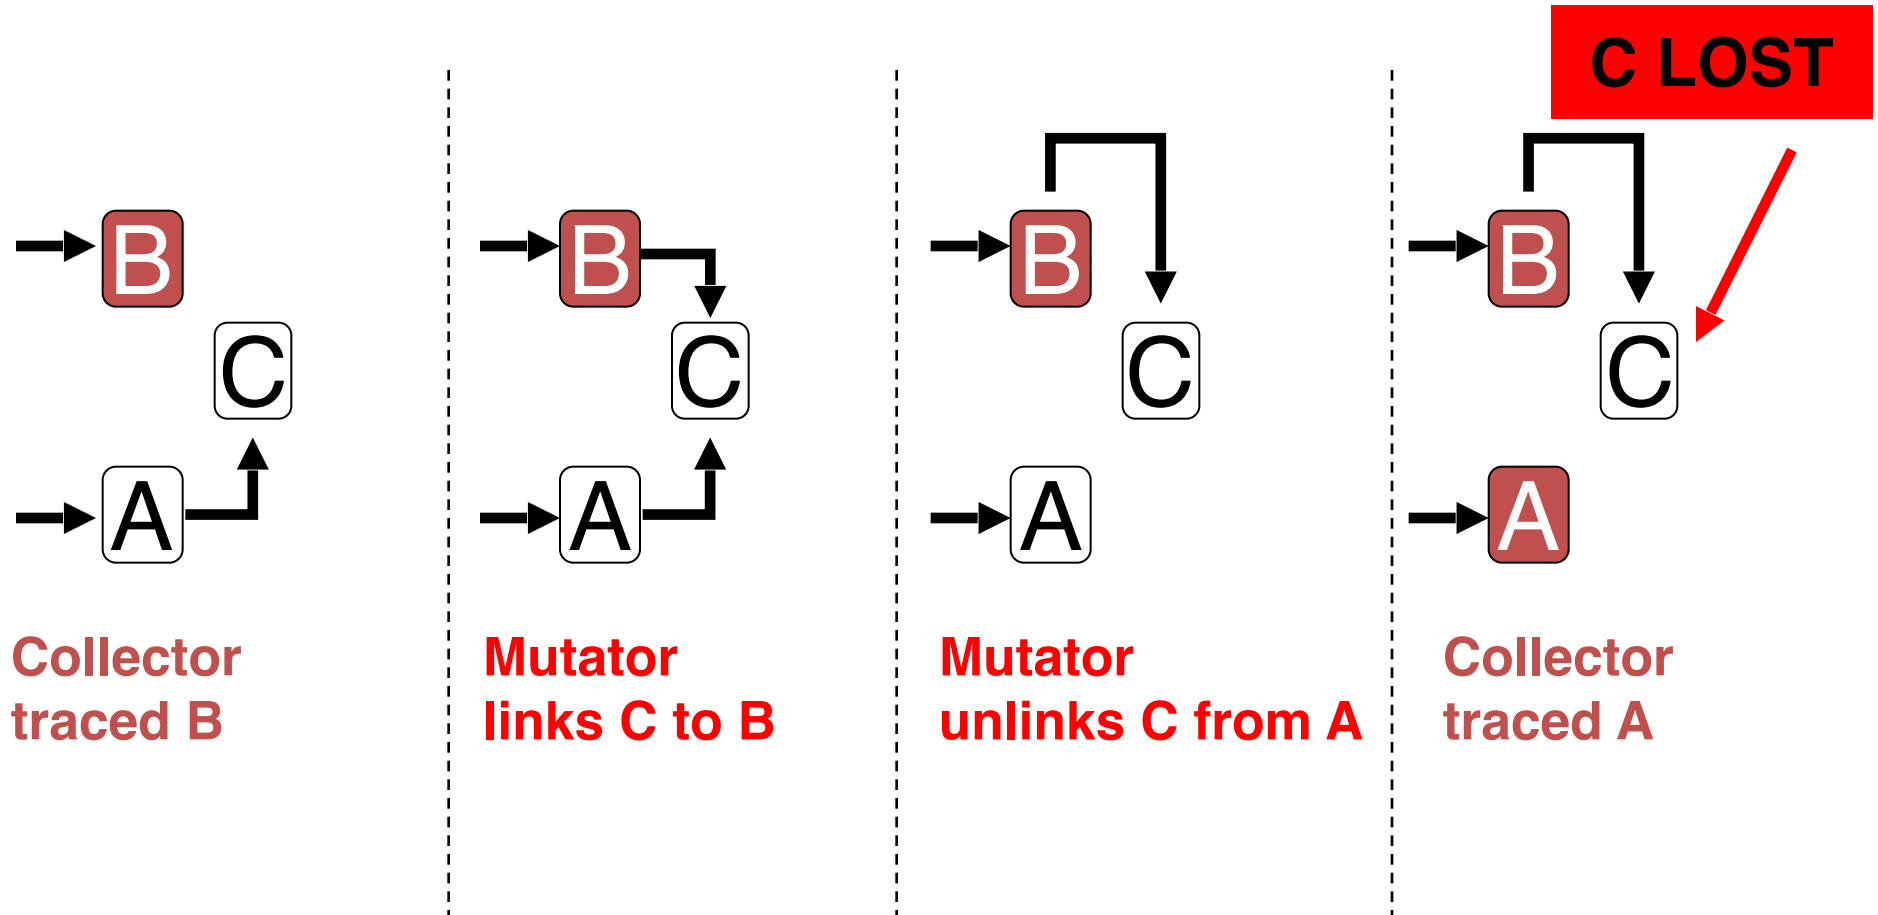
Concurrent GC: interference

SYSTEM = MUTATOR II COLLECTOR



Concurrent GC: interference


SYSTEM = MUTATOR II COLLECTOR



Concurrent GC: families of algorithms


DIJKSTRA

Marks C when
C is linked to B




YUASA

Marks C when
link to C is removed



STEELE


Rescan B when
C is linked to B



Concurrent GC: families of algorithms


DIJKSTRA

Marks C when
C is linked to B




YUASA

Marks C when
link to C is removed



STEELE


Rescan B when
C is linked to B



Concurrent GC: families of algorithms


DIJKSTRA

Marks C when
C is linked to B




YUASA

Marks C when
link to C is removed



STEELE


Rescan B when
C is linked to B



Concurrent GC: families of algorithms


DIJKSTRA

Marks C when
C is linked to B




YUASA

Marks C when
link to C is removed



STEELE


Rescan B when
C is linked to B




Copying GC

- Partition the heap into two halves
 - old space
 - new space
- Copy all reachable objects from old space to new space
- Swap roles of old and new spaces


Example: Copying GC



Example: Copying GC



Example: Copying GC



Properties of Copying GC

- Major disadvantage: **half of the heap is not used**
- Compaction for free
- Touch only the live objects
 - **good when most objects are dead**
 - usually most new objects are dead
- **Generational GC:** use a small space for **young** objects and collect this space using copying GC

Very simplistic comparison

	Reference Counting	Mark and Sweep	Copying
Complexity	Pointer updates + dead objects	Size of heap (live objects)	Live objects
Space overhead	Count/object + stack for DFS	Bit/object + stack for DFS	Half heap wasted
Compaction	Additional work	Additional work	For free
Pause time	Mostly short	long	long
More issues	Cycle collection		

Modern memory management

- Considers standard program properties
- Parallelism
 - stop the program and collect in parallel on all available processors
 - run collection concurrently with the program run
- Cache consciousness
- Real-time

Conservative GC

- Any value can be cast down to a pointer in C
- How can we follow pointers in a structure?

Conservative GC

- Any value can be cast down to a pointer in C
- How can we follow pointers in a structure?
- Easy: conservatively consider anything that can be a pointer to be a pointer
- Practical! Boehm collector

Conservative GC

- Any value can be cast down to a pointer in C
- How can we follow pointers in a structure?
- Easy: conservatively consider anything that can be a pointer to be a pointer
- Practical! Boehm collector
- Can we implement a conservative copying GC?

Conservative GC

- Any value can be cast down to a pointer in C
- How can we follow pointers in a structure?
- Easy: conservatively consider anything that can be a pointer to be a pointer
- Practical! Boehm collector
- Can we implement a conservative copying GC?
- No. Cannot update pointers to the new address: if we don't know whether the value is a pointer, we cannot update it

Summary: Garbage Collection

- Reference counting
- Mark and sweep
- Compaction
- Copying
- Generational
- Parallel
- Concurrent

ERROR HANDLING

Runtime checks

- Generate code for checking attempted illegal operations
 - null pointer check
 - array length, virtual call, reference arguments to library call
 - array bounds check
 - array allocation size check
 - division by zero
 - ...
- If check fails jump to **error handler** code that prints a message and gracefully exits program
- Alternatively, use an **exception handling** mechanism

Null pointer check

```
# null pointer check  
cmp $0,%eax  
je labelNPE
```

Single generated handler for entire program

```
labelNPE:  
    push $strNPE          # error message  
    call __println  
    push $1               # error code  
    call __exit
```

Array bounds check

```
# array bounds check
mov -4(%eax),%ebx    # ebx = length
mov $0,%ecx          # ecx = index
cmp %ecx,%ebx
jle labelABE         # ebx <= ecx ?
cmp $0,%ecx
jl  labelABE         # ecx < 0 ?
```

Single generated handler for entire program

```
labelABE:
    push $strABE      # error message
    call __println
    push $1           # error code
    call __exit
```


Array allocation size check

```
# array size check  
cmp $0,%eax      # eax == array size  
jle labelASE     # eax <= 0 ?
```

Single generated handler for entire program

```
labelASE:  
    push $strASE    # error message  
    call __println  
    push $1         # error code  
    call __exit
```

Exceptions



Exception example

```
org.eclipse.swt.SWTException: Graphic is disposed
    at org.eclipse.swt.SWT.error(SWT.java:3744)
    at org.eclipse.swt.SWT.error(SWT.java:3662)
    at org.eclipse.swt.SWT.error(SWT.java:3633)
    at org.eclipse.swt.graphics.GC.getClipping(GC.java:2266)
    at com.aELITIS.azureus.ui.swt.views.list.ListRow.doPaint(ListRow.java:260)
    at com.aELITIS.azureus.ui.swt.views.list.ListRow.doPaint(ListRow.java:237)
    at com.aELITIS.azureus.ui.swt.views.list.ListView.handleResize(ListView.java:867)
    at com.aELITIS.azureus.ui.swt.views.list.ListView$5$2.runSupport(ListView.java:406)
    at org.gudy.azureus2.core3.util.AERunnable.run(AERunnable.java:38)
    at org.eclipse.swt.widgets.RunnableLock.run(RunnableLock.java:35)
    at org.eclipse.swt.widgets.Synchronizer.runAsyncMessages(Synchronizer.java:130)
    at org.eclipse.swt.widgets.Display.runAsyncMessages(Display.java:3323)
    at org.eclipse.swt.widgets.Display.readAndDispatch(Display.java:2985)
    at org.gudy.azureus2.ui.swt.mainwindow.SWTThread.<init>(SWTThread.java:183)
    at org.gudy.azureus2.ui.swt.mainwindow.SWTThread.createInstance(SWTThread.java:67)
```

....